

Tech-Quantum

Microsoft Dynamics | AI | ML | More...

[HOME](#)[ARTIFICIAL INTELLIGENCE \(AI\)](#)[QUANTUM COMPUTING](#)[DYNAMICS 365](#)[MICROSOFT AZURE](#)[ECOSYSTEM](#)[TRAINING VIDEOS](#)[AUTHORS](#)

[TECH-QUANTUM](#) > [AI & ML](#) > [DEEP LEARNING](#) > [IMPLEMENT BACK PROPAGATION IN NEURAL NETWORKS](#)

Implement Back Propagation in Neural Networks

Posted on October 6, 2018 by Deepak Battini

When building neural networks, there are several steps to take. Perhaps the two most important steps are implementing forward and backward propagation. Both these terms sound really heavy and are always scary to beginners. The absolute truth is that these techniques can be properly understood if they are broken down into their individual steps. In this tutorial, we will focus on backpropagation and the intuition behind every step of it.

What is Back Propagation?

This is simply a technique in implementing neural networks that allow us to calculate the gradient of parameters in order to perform gradient descent and minimize our cost function. Numerous scholars have described back propagation as arguably the most mathematically intensive part of a neural network. Relax though, as we will completely decipher every part of back propagation in this tutorial.



SUBMIT YOUR VIDEO REQUEST

[Click Here](#)




RECENT POSTS

[How to work with custom help panes and guided tasks?](#)

[How to work with Global Custom Help URL & Table Help URL?](#)

[How to add sample data in Dynamics 365 Customer Engagement?](#)

Backpropagation Learning



Backpropagation Learning

$$E_{out\ i} = d_{out\ i} - out_i$$

$$E_{total} = \sum_{i=0}^{num(n_{out})} E_{out\ i}^2$$

$$E_{hid\ i} = \sum_{k=1}^{num(n_{out})} E_{out\ k} \cdot w_{out\ i,k}$$

$$diff_{hid\ i} = E_{hid\ i} \cdot (1 - o(n_{hid\ i})) \cdot o(n_{hid\ i})$$

10

Implementing Back Propagation

Assuming a simple two-layer neural network – one hidden layer and one output layer. We can perform back propagation as follows

Initialize the weight and bias to be used for the neural network:

This involves randomly initializing the weights and biases of the neural networks. The gradient of these parameters will be obtained from the backward propagation and used to update gradient descent.

```
1. #Import Numpy library
2. import numpy as np
3.
4. #set seed for reproducibility
5. np.random.seed(100)
```

```
1. #We will first initialize the weights and
   bias needed and store them in a dictionary
   called W_B
2. def initialize(num_f, num_h, num_out):
3.
4.     '''
5.     Description: This function randomly
   initializes the weights and biases of each
   layer of the neural network
6.
7.     Input Arguments:
8.     num_f - number of training features
```

Many world interpretations for building a powerful computer—Is wave the reality?

Many world interpretation towards building a powerful computer

TAGS

365 2016 ai artificial intelligence artificial neural network AZ900 azure business business central CDS central cloud CommonDataService computer vision computing crm csharp customer service d65 D365 data deep learning dynamics dynamics 365 dynamics365 fundamentals how to howto learning machine machine learning manipulation microsoft microsoft dynamics 365 ml neural network opencv PowerApps

[project service](#)
[automation psa](#)
[SDK service training](#)
[XRM YouTube](#)

```
9.         num_h -the number of nodes in the hidden
           layers
10.         num_out - the number of nodes in the
           output
11.
12.         Output:
13.
14.         W_B - A dictionary of the initialized
           parameters.
15.
16.         '''
17.
18.         #randomly initialize weights and biases,
           and proceed to store in a dictionary
19.         W_B = {
20.             'W1': np.random.randn(num_h, num_f),
21.             'b1': np.zeros((num_h, 1)),
22.             'W2': np.random.randn(num_out,
           num_h),
23.             'b2': np.zeros((num_out, 1))
24.         }
25.         return W_B
```

Perform forward propagation: This involves calculating both the linear and activation outputs for both the hidden layer and the output layer.

For the hidden layer:

We will be using the relu activation function as seen below:

```
1.         #We will now proceed to create functions for
           each of our activation functions
2.
3.         def relu (Z):
4.
5.             '''
6.             Description: This function performs the
           relu activation function on a given number or
           matrix.
7.
8.             Input Arguments:
9.             Z - matrix or integer
10.
11.             Output:
```

```

12.
13.     relu_Z - matrix or integer with relu
        performed on it
14.
15.     '''
16.     relu_Z = np.maximum(Z,0)
17.
18.     return relu_Z

```

For the Output layer:

We will be using the sigmoid activation function as seen below:

```

1.  def sigmoid (Z):
2.
3.      '''
4.      Description: This function performs the
        sigmoid activation function on a given number
        or matrix.
5.
6.      Input Arguments:
7.      Z - matrix or integer
8.
9.      Output:
10.
11.      sigmoid_Z - matrix or integer with
        sigmoid performed on it
12.
13.      '''
14.      sigmoid_Z = 1 / (1 + (np.exp(-Z)))
15.
16.      return sigmoid_Z

```

Perform the forward propagation:

```

1.  #We will now proceed to perform forward
    propagation
2.
3.  def forward_propagation(X, W_B):
4.      '''
5.      Description: This function performs the
        forward propagation in a vectorized form
6.
7.      Input Arguments:
8.      X - input training examples
9.      W_B - initialized weights and biases
10.

```

```

11.         Output:
12.
13.         forward_results - A dictionary containing
the linear and activation outputs
14.
15.         '''
16.
17.         #Calculate the linear Z for the hidden
layer
18.         Z1 = np.dot(X, W_B['W1']).T + W_B['b1']
19.
20.         #Calculate the activation ouput for the
hidden layer
21.         A = relu(Z1)
22.
23.         #Calculate the linear Z for the output
layer
24.         Z2 = np.dot(A, W_B['W2']).T + W_B['b2']
25.
26.         #Calculate the activation ouput for the
ouptu layer
27.         Y_pred = sigmoid(Z2)
28.
29.         #Save all ina dictionary
30.         forward_results = {"Z1": Z1,
31.                             "A": A,
32.                             "Z2": Z2,
33.                             "Y_pred": Y_pred}
34.
35.         return forward_results

```

Perform Backward Propagation: Calculate the gradients of the cost relative to the parameters relevant for gradient descent. In this case, $dLdZ_2$, $dLdW_2$, $dLdb_2$, $dLdZ_1$, $dLdW_1$ and $dLdb_1$. These parameters will be combined with the learning rate to perform gradient descent. We will implement a vectorized version of the backpropagation for a number of training samples – no_examples.

The step by step guide is as follows:

- **Obtain Results from forwarding propagation as seen below:**

```

1. forward_results = forward_propagation(X, W_B)
2. Z1 = forward_results['Z1']
3. A = forward_results['A']
4. Z2 = forward_results['Z2']
5. Y_pred = forward_results['Y_pred']

```

- **Obtain the number of training samples as seen below:**

```

1. no_examples = X.shape[1]

```

- **Calculate the Loss of the function:**

```

1. L = (1/no_examples) * np.sum(-Y_true *
    np.log(Y_pred) - (1 - Y_true) * np.log(1 -
    Y_pred))

```

- **Calculate the gradient for each parameter as seen below:**

```

1. dLdZ2= Y_pred - Y_true
2. dLdW2 = (1/no_examples) * np.dot(dLdZ2, A.T)
3. dLdb2 = (1/no_examples) * np.sum(dLdZ2,
    axis=1, keepdims=True)
4. dLdZ1 = np.multiply(np.dot(W_B['W2'].T,
    dLdZ2), (1 - np.power(A, 2)))
5. dLdW1 = (1/no_examples) * np.dot(dLdZ1, X.T)
6. dLdb1 = (1/no_examples) * np.sum(dLdZ1,
    axis=1, keepdims=True)

```

- **Store the calculated gradients needed for gradient descent in a dictionary:**

```

1. gradients = {"dLdW1": dLdW1,
2.              "dLdb1": dLdb1,
3.              "dLdW2": dLdW2,
4.              "dLdb2": dLdb2}

```

- **Return the loss and the stored gradients:**

```

1. return gradients, L

```

Here is the complete function:

```

1. def backward_propagation(X, W_B, Y_true):
2.     '''Description: This function performs
    the backward propagation in a vectorized form

```

```

3.
4.     Input Arguments:
5.     X - input training examples
6.     W_B - initialized weights and biases
7.     Y_True - the true target values of the
training examples
8.
9.     Output:
10.
11.     gradients - the calculated gradients of
each parameter
12.     L - the loss function
13.
14.     '''
15.
16.     # Obtain the forward results from the
forward propagation
17.
18.     forward_results = forward_propagation(X,
W_B)
19.     Z1 = forward_results['Z1']
20.     A = forward_results['A']
21.     Z2 = forward_results['Z2']
22.     Y_pred = forward_results['Y_pred']
23.
24.     #Obtain the number of training samples
25.     no_examples = X.shape[1]
26.
27.     # Calculate loss
28.     L = (1/no_examples) * np.sum(-Y_true *
np.log(Y_pred) - (1 - Y_true) * np.log(1 -
Y_pred))
29.
30.     #Calculate the gradients of each
parameter needed for gradient descent
31.     dLdZ2= Y_pred - Y_true
32.     dLdW2 = (1/no_examples) * np.dot(dLdZ2,
A.T)
33.     dLdb2 = (1/no_examples) * np.sum(dLdZ2,
axis=1, keepdims=True)
34.     dLdZ1 = np.multiply(np.dot(W_B['W2'].T,
dLdZ2), (1 - np.power(A, 2)))
35.     dLdW1 = (1/no_examples) * np.dot(dLdZ1,
X.T)
36.     dLdb1 = (1/no_examples) * np.sum(dLdZ1,
axis=1, keepdims=True)
37.

```

```
38.         #Store gradients for gradient descent in
          a dictionary
39.         gradients = {"dLdW1": dLdW1,
40.                      "dLdb1": dLdb1,
41.                      "dLdW2": dLdW2,
42.                      "dLdb2": dLdb2}
43.
44.         return gradients, L
```

Many people always think that back propagation is difficult, but as you have seen in this tutorial, it is not. Understanding every step is imperative to grasp the entire back-propagation technique. Also, it pays to have a good grasp of mathematics – linear algebra and calculus – in order to understand how the individual gradients of each function are calculated. With these tools, back propagation should be a piece of cake! In practice, back propagation is usually handled for you by the Deep Learning framework you are using. However, it pays to understand the inner workings of this technique as it can sometimes help you understand why your neural network may not be training well.

Here is the link to jupyter [notebook](#)

I will post more blogs in future about back propagation to get you deeper understanding, so stay tuned:)

Posted in AI & ML, Deep Learning Tagged artificial intelligence, backpropagation, machine learning, python

[← Reinforcement learning – How machine can learn new skills?](#)

[Awesome seaborn for Data Visualization – Part 2 →](#)

Leave a Reply

Your email address will not be published. Required fields are marked *

Comment *

Name *

Email *

Website

Post Comment

Theme by Out the Box

[Contact Us](#) | [About Authors](#) | [SiaNet – Resources](#) | [__Blank__](#) | [Tech-Quantum Ecosystem](#) | [Videos](#)